

Amazon Web Services (AWS) en tiempos de Big Data



Dr. Andrea Villanes

Materials:

https://github.com/andreavillanes/AWS_EMR

Un poco sobre mi...

- PhD en Computer Science **NC STATE UNIVERSITY**

- Assistant Professor 



- [Master in Science of Analytics](#) **NC STATE UNIVERSITY**

- Gracias DMC por la organizacion!







Services ▾

Resource Groups ▾




EC2



S3



RDS



EMR



Andrea-AWS ▾

N. California ▾

Support ▾

Amazon EMR

Clusters

Security configurations

Block public access

VPC subnets

Events

Notebooks

Git repositories

Help

What's new

Notebook: Spark_4_22 Ready Notebook is ready to run jobs on cluster j-210FMVO29M0YF.

Open in JupyterLab

Open in Jupyter

Stop

Delete

Notebook



Notebook ID: e-ENEY4TDJ8F6WY1MJ6KEO1YW14

Description: --

Last modified: 1 minute ago ⓘ

Last modified by: ...root ⓘ

Created on: 2020-04-22 13:50 (UTC-4)

Created by: ...root ⓘ

Service IAM role: [EMR_Notebooks_DefaultRole](#) ↗

Notebook tags: creatorUserId = 799535928231 [View All](#) / [Edit](#)

Notebook location: s3://aws-emr-resources-799535928231-us-west-1/notebooks/ 📁

Cluster

Cluster: NotebookCluster

Cluster Id: [j-210FMVO29M0YF](#)

Cluster status: Waiting Cluster ready to run steps.

Cluster tags: creator = NOTEBOOK_CONSOLE [View All](#)

Step logs: s3://aws-logs-799535928231-us-west-1/elasticmapreduce/ 📁

Git repositories



The repository can be linked to a notebook once the notebook is ready. Make sure your cluster, service role and security groups have the required settings. [Learn more](#) ↗



Services ▾

Resource Groups ▾




EC2



S3



RDS



EMR



Andrea-AWS ▾

N. California ▾

Support ▾

Amazon EMR

Clusters

Security configurations

Block public access

VPC subnets

Events

Notebooks

Git repositories

Help

What's new

Notebook: Spark_4_22 Ready Notebook is ready to run jobs on cluster j-210FMVO29M0YF.

Open in JupyterLab

Open in Jupyter

Stop

Delete

Notebook



Notebook ID: e-ENEY4TDJ8F6WY1MJ6KEO1YW14

Description: --

Last modified: 1 minute ago ⓘ

Last modified by: ...root ⓘ

Created on: 2020-04-22 13:50 (UTC-4)

Created by: ...root ⓘ

Service IAM role: [EMR_Notebooks_DefaultRole](#) ↗

Notebook tags: creatorUserId = 799535928231 [View All](#) / [Edit](#)

Notebook location: s3://aws-emr-resources-799535928231-us-west-1/notebooks/ 📁

Cluster

Cluster: NotebookCluster

Cluster Id: [j-210FMVO29M0YF](#)

Cluster status: Waiting Cluster ready to run steps.

Cluster tags: creator = NOTEBOOK_CONSOLE [View All](#)




Step logs: s3://aws-logs-799535928231-us-west-1/elasticmapreduce/ 📁

Git repositories






The repository can be linked to a notebook once the notebook is ready. Make sure your cluster, service role and security groups have the required settings. [Learn more](#) ↗





Notebook



Console




Other




+		+	↑	↺	⚙
/					
Name		Last Modified			
Spark_4_22.ipynb		9 minutes ago			

Launcher


Notebook



Console



Other



Agenda

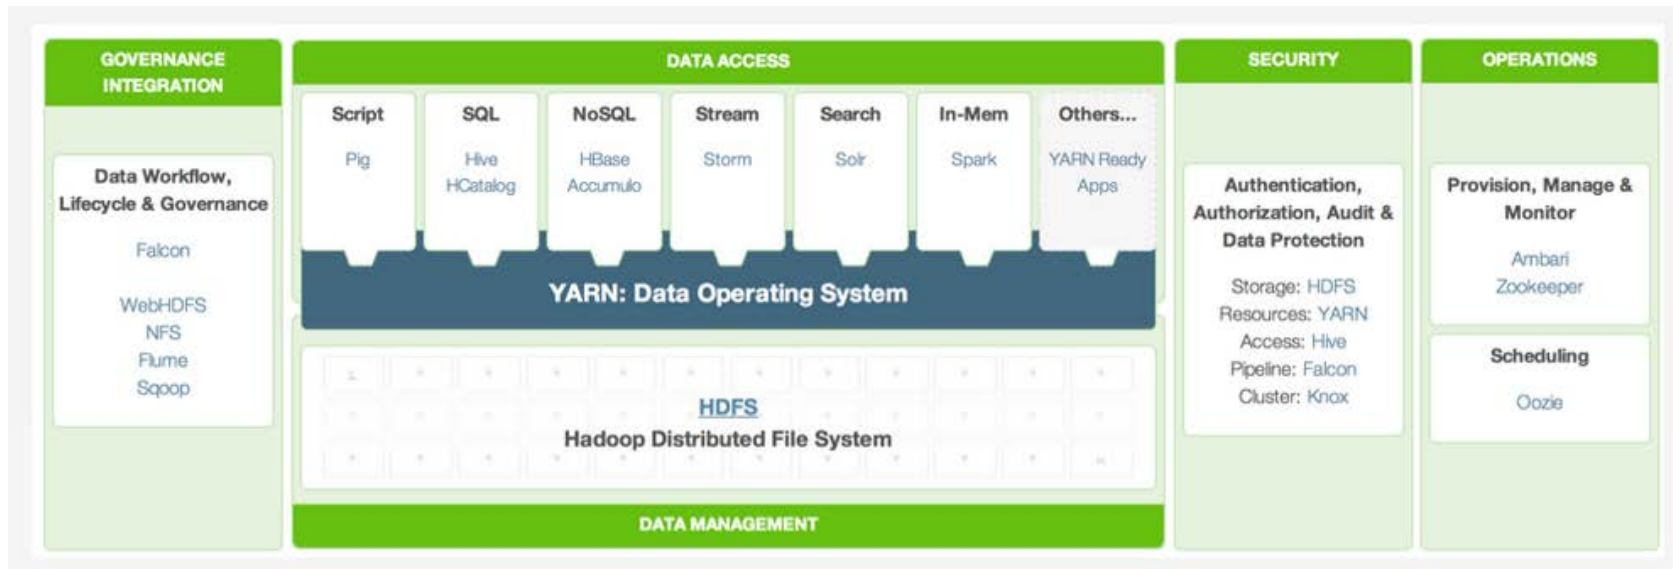
- Apache Hadoop
- Apache Spark
- Distribuciones Comerciales de Hadoop
- AWS Elastic Map Reduce (EMR)
- PySpark
- Spark SQL
- Spark MLlib
- Como abrir una cuenta de AWS gratis?

Motivacion para Hadoop



- Velocidad
- Variedad
- Volumen
- Data tiene valor
- Dos problemas que tenemos que resolver:
 1. Como podemos almacenar grandes cantidades de datos a un costo razonable?
 2. Como podemos analizar la data que hemos almacenado?

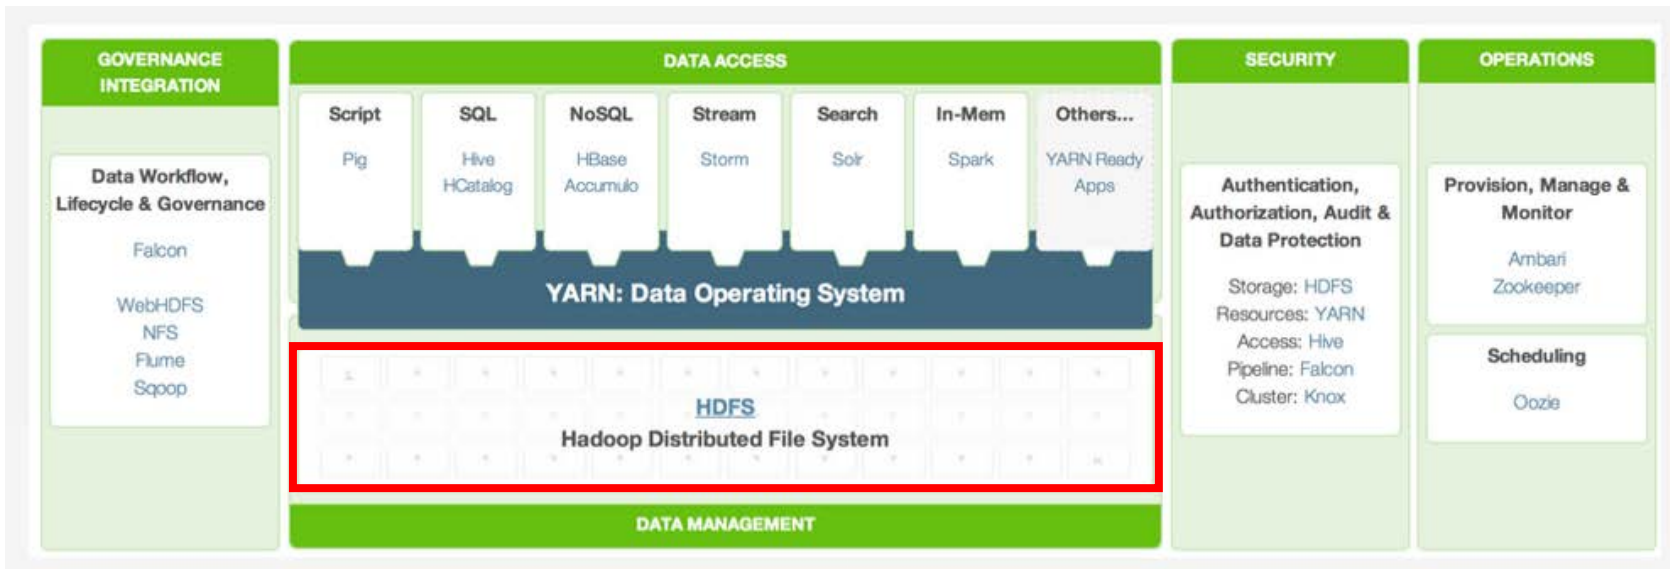
Apache Hadoop



Beneficios:

1. Escalabilidad
2. Tolerancia a las fallas

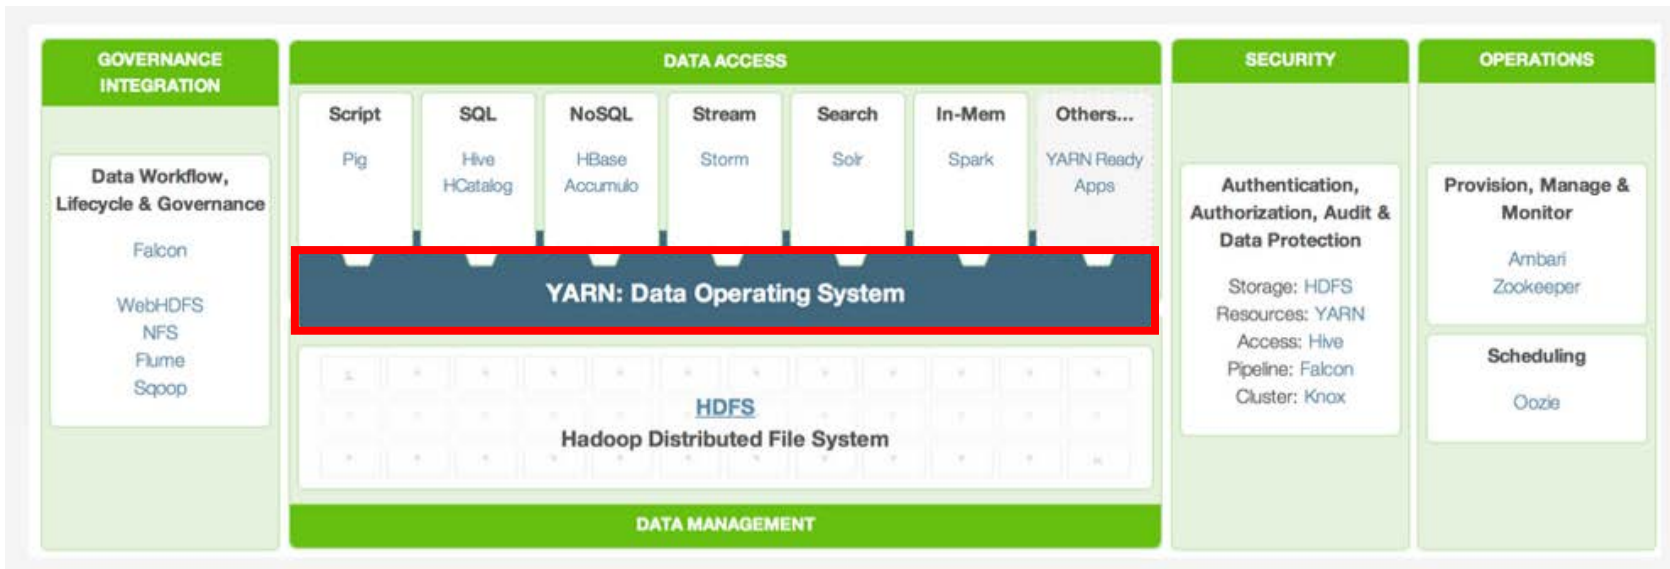
Apache Hadoop



Beneficios:

1. Escalabilidad
2. Tolerancia a las fallas

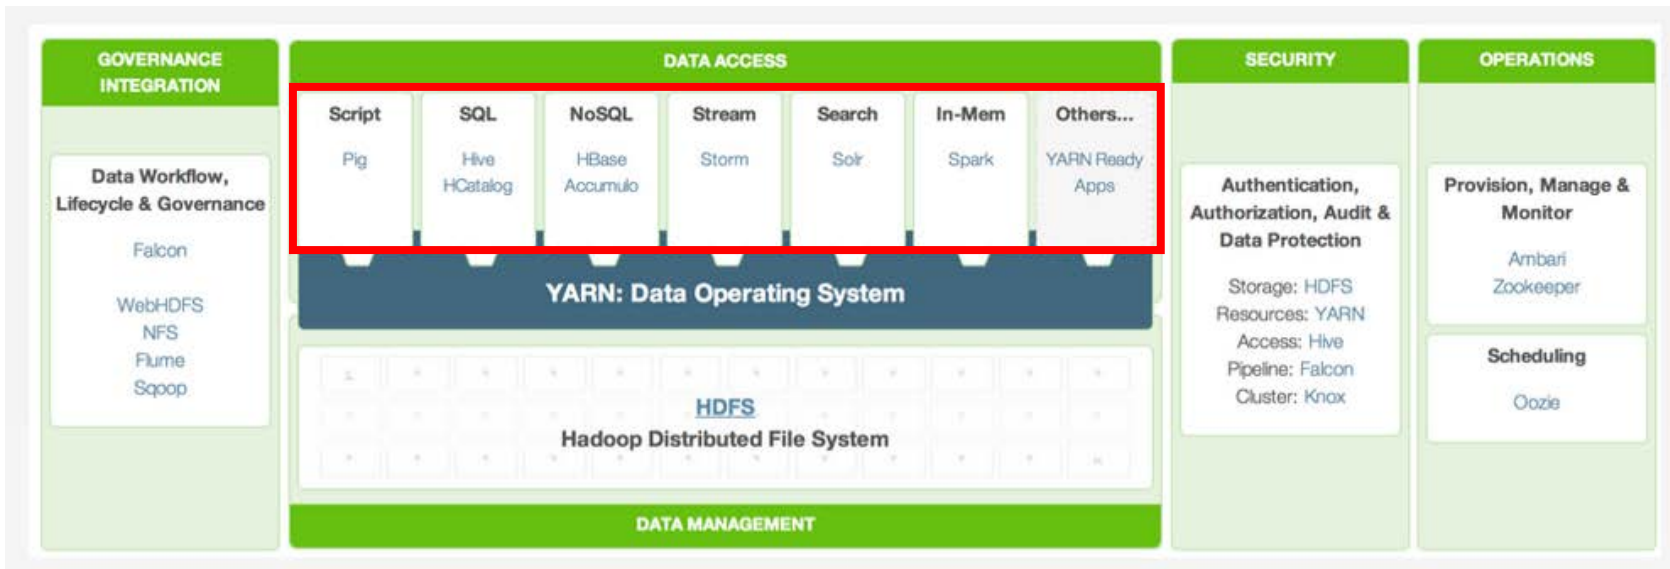
Apache Hadoop



Beneficios:

1. Escalabilidad
2. Tolerancia a las fallas

Apache Hadoop



Beneficios:

1. Escalabilidad
2. Tolerancia a las fallas


Motivacion para Spark



- Focus en operaciones interactivas e iterativas
- Utiliza procesamiento en memoria
- Los hace ideal para aplicaciones de data science (data mining, machine learning)
- Extensive API support for Java, Scala, R and Python
- **Runs Everywhere:** Spark runs on Hadoop, standalone, or in the cloud. It can access diverse data sources including HDFS, Cassandra, HBase, and AWS S3.

Apache Spark Components

- Combine SQL, streaming, and complex analytics.





Preguntas?

Agenda

- Apache Hadoop
- Apache Spark
- Distribuciones Comerciales de Hadoop
- AWS Elastic Map Reduce (EMR)
- PySpark
- Spark SQL
- Spark MLlib
- Como abrir una cuenta de AWS gratis?

Apache Hadoop y Apache Spark son proyectos open source



open source



Necesitamos...



Necesitamos que los requerimientos para Hadoop estén alineados con las necesidades de las empresas, y eso resulta en el nacimiento de distribuciones comerciales.

Ejemplos de Hadoop vendors que venden una Distribucion de Hadoop



Distribuciones de Hadoop



Hortonworks Data Platform (HDP)



Cloudera CDH



MapR Distribution



Microsoft's Azure HDInsight

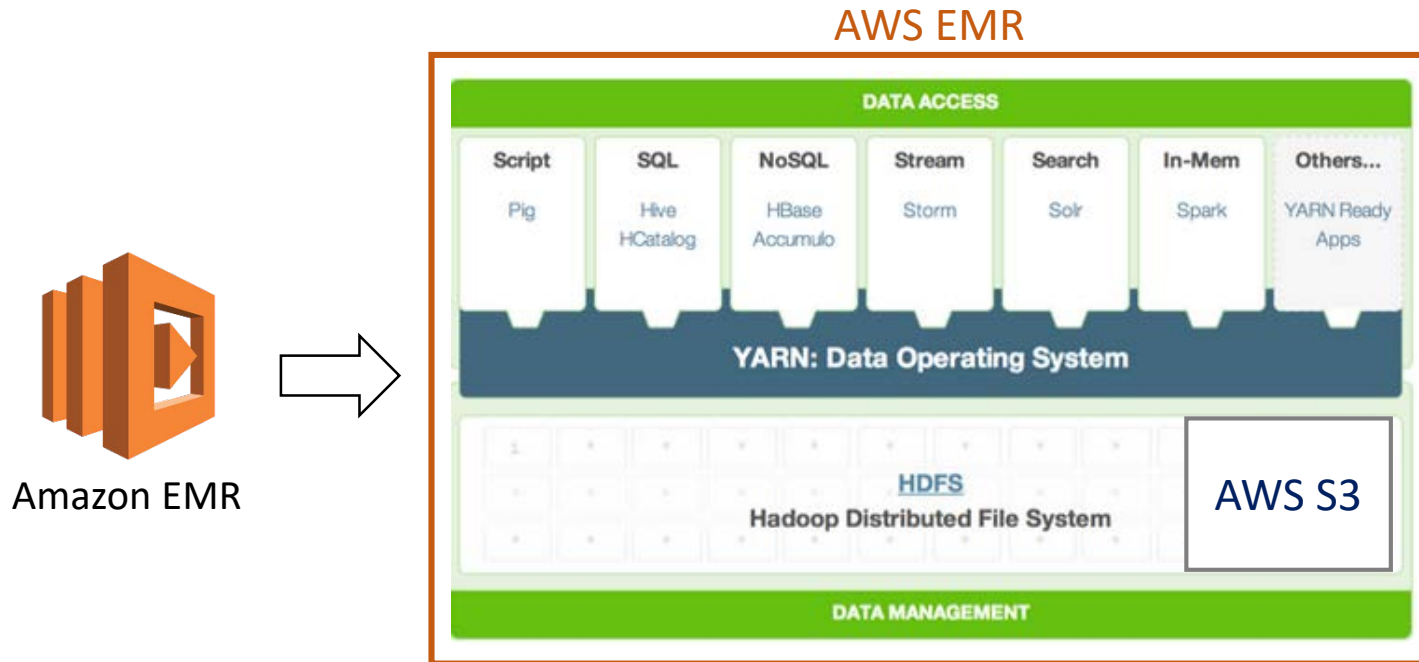


Amazon Elastic MapReduce (EMR)



Cloud Dataproc

AWS Hadoop Distribution: Amazon Elastic MapReduce (EMR)



- Amazon EMR includes EMRFS, a connector allowing Hadoop to use **S3 as a storage layer**.
- **HDFS** is automatically installed with Hadoop on your EMR cluster, and **you can use HDFS along with Amazon S3 to store your input and output data**.
- Amazon EMR configures Hadoop to use **HDFS for intermediate data** created during MapReduce jobs, even if your input data is located in Amazon S3.

Amazon EMR programmatically installs and configures applications in the Hadoop project, including Hadoop MapReduce (YARN), and HDFS, across the nodes in your cluster.

Amazon Simple Storage Service (S3)



- S3: **object storage service**.
- Objects are stored in **buckets**.
- Natively online, **HTTP access**.
- Every object in Amazon S3 can be **uniquely addressed through the combination of the web service endpoint, bucket name, key, and optionally, a version**.
- **Store and retrieve any amount of data**, any time, from anywhere on the web.
- Amazon S3 is **highly scalable, reliable, low cost, and designed for durability**.
- Data **can be stored as-is**: there is no need to convert it to a predefined schema.

Preguntas?

Agenda

- Apache Hadoop
- Apache Spark
- Distribuciones Comerciales de Hadoop
- AWS Elastic Map Reduce (EMR)
- PySpark
- Spark SQL
- Spark MLlib
- Como abrir una cuenta de AWS gratis?

PySpark

Python + Spark = PySpark

- PySpark is the collaboration of Apache Spark and Python.
- **Apache Spark** is an open-source cluster-computing framework, built around speed, ease of use, and streaming analytics.
- **Python** is a general-purpose, high-level programming language.



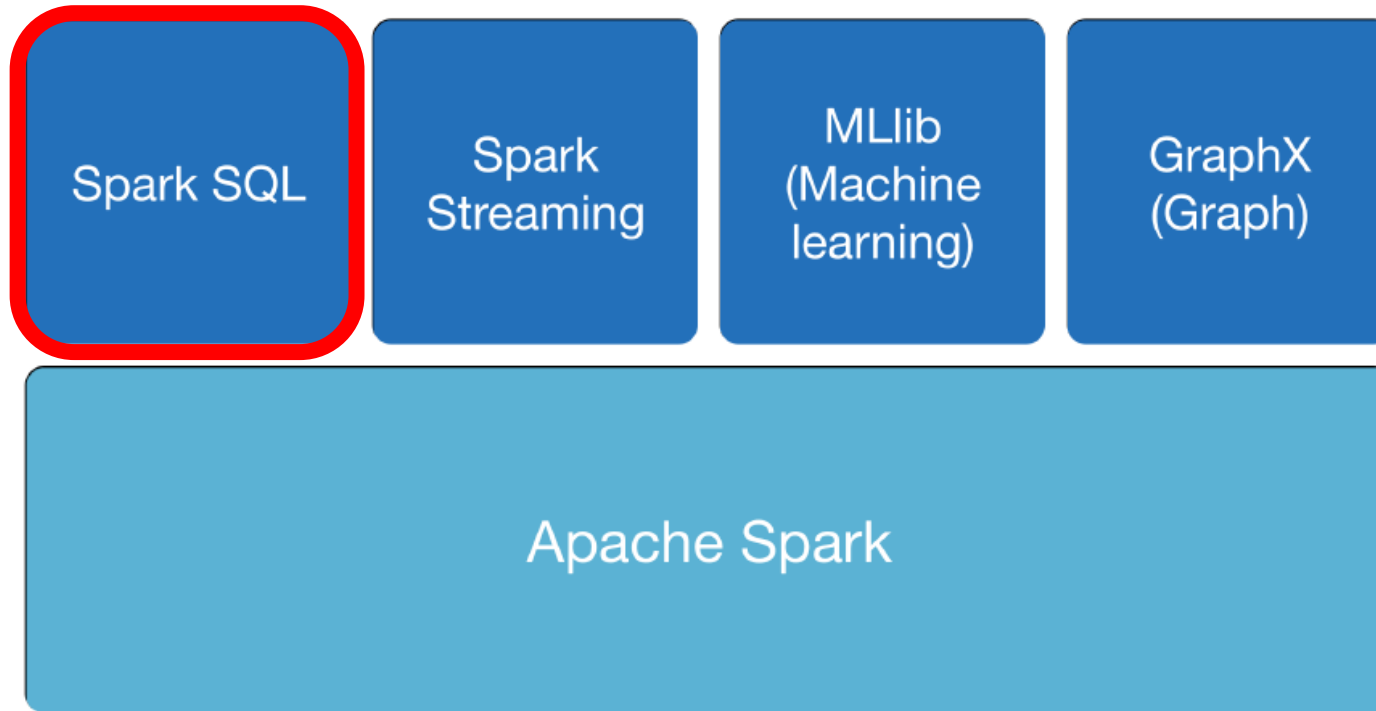
```
df = spark.read.json("logs.json")
df.where("age > 21")
  .select("name.first").show()
```

Spark's Python DataFrame API

Spark SQL

Apache Spark Components

- Combine SQL, streaming, and complex analytics.



Apache Spark SQL



- Spark SQL is a Spark's module for working with **structured data**.
- There are several ways to interact with Spark SQL including SQL, the DataFrames API and the Datasets API. However, PySpark only implements the DataFrames API.
- A DataFrame is a distributed collection of data organized into named columns. It is **conceptually equivalent to a table in a relational database or a data frame in R/Python**, but with richer optimizations under the hood. DataFrames can be constructed from a wide array of [sources](#) such as: structured data files, tables in Hive, external databases, or existing RDDs.
- Executes SQL queries written using either a basic SQL syntax

PySpark & Spark SQL Cheat Sheet

Spark SQL is Apache Spark's module for working with structured data.



Initializing SparkSession

A **SparkSession** can be used to create **DataFrame**, register **DataFrame** as views, execute SQL over views, and read csv, json, txt and parquet files.

In EMR Notebooks, **SparkSession** is automatically created for you. The **SparkSession** is accessible through a variable called *spark*. If you want to verify your Spark version:

```
>>> spark.version
```

Creating DataFrames

From Spark Data Sources

CSV

```
>>> df = spark.read.csv("s3a://bucket_name/airlines.csv", inferSchema = True, header=True)
```

JSON

```
>>> df2 = spark.read.json("s3a://bucket_name/customer.json")
```

```
>>> df2.show()
```

```
>>> df3 = spark.read.load("s3a://bucket_name/people.json", format="json")
```

Parquet files

```
>>> df4 = spark.read.load("s3a://bucket_name/users.parquet")
```

TXT files

```
>>> df5 = spark.read.text("s3a://bucket_name/people.txt")
```

View the DataFrame

Show() - Displays the top 20 rows of **DataFrame** in a tabular form.

```
>>> df.show()
```

Show(n) - Displays the top *n* rows of **DataFrame** in a tabular form.

```
>>> df.show(n)
```

Inspect Data

```
>>> df.describe().show()
```

```
>>> df.columns
```

```
>>> df.count()
```

```
>>> df.distinct().count()
```

```
>>> df.printSchema()
```

Compute summary statistics

Return the columns of **df**

Count the number of rows in **df**

Count the number of distinct rows in **df**

Print the schema of **df**

Queries

```
>>> from pyspark.sql import functions as F
Select
>>> df.select("firstName").show()
>>> df.select("firstName", "lastName") \
    .show()
>>> df.select(df["firstName"], df["age"] + 1)
>>> df.select(df["age"] > 24).show()
```

Show all entries in *firstName* column
Show all entries in *firstName*, and *lastName*
Show all entries in *firstName* and *age*, add 1 to the entries of *age*
Show all entries where *age* > 24

When

```
>>> df.select("firstName",
    F.when(df.age > 30, 1) \
    .otherwise(0)) \
    .show()
>>> df[df.firstName.isin("Jane", "Boris")]
    .show()
```

Show *firstName* and 0 or 1 depending on *age* > 30

Show *firstName* if in the given options

Like

```
>>> df.select("firstName",
    df.lastName.like("Smith")) \
    .show()
```

Show *firstName*, and *lastName* if *lastName* is like Smith

Startswith – Endswith

```
>>> df.select("firstName",
    df.lastName \
    .startswith("Sm")) \
    .show()
>>> df.select(df.lastName.endswith("th")) \
    .show()
```

Show *firstName*, and *lastName* if *lastName* starts with *Sm*

Show last names ending in *th*

Substring

```
>>> df.select(df.firstName.substr(1, 3) \
    .alias("name")) \
    .collect()
```

Return substrings of *firstName*

Between

```
>>> df.select(df.age.between(22, 24)) \
    .show()
```

Show *age* if values between 22 and 24

Adding Columns

```
>>> from pyspark.sql.functions import log
```

```
>>> df2 = df.withColumn("new_column", log("rating"))
```

```
>>> df2.show()
```

Duplicate Values

```
>>> df = df.dropDuplicates()
```

GroupBy

```
>>> df.groupBy("age") \
    .count() \
    .show()
```

Group by *age*, count the members in the groups

Filter

```
>>> df.filter(df["age"] > 24).show()
```

Filter entries of *age*, only keep those records of which the values are > 24

Missing & Replacing Values

```
>>> df.na.fill(50).show()
>>> df.na.drop().show()
>>> df.na \
    .replace(10, 20) \
    .show()
```

Replace null values
Return new **df** omitting rows with null values
Return new **df** replacing one value with another

Running SQL Queries Programmatically

Registering DataFrames as Views

```
>>> df.createOrReplaceTempView("customer")
```

Query Views

```
>>> sqlDF = spark.sql("SELECT * FROM customer").show()
```

Output

Data Structure

```
>>> rdd1 = df.rdd
>>> df.toPandas()
```

Convert *df* into an **RDD**
Return the contents of *df* as **Pandas DataFrame**

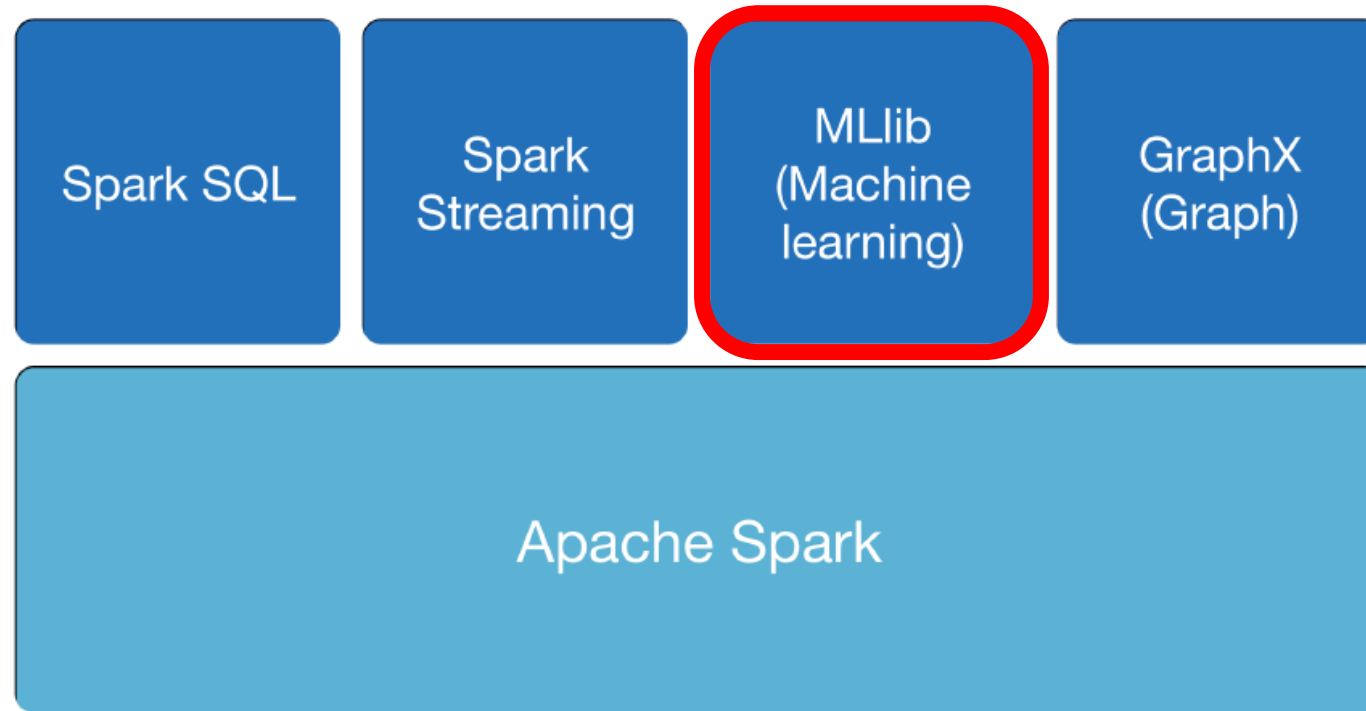
Write & Save to Files

```
>>> df.select("firstName", "city") .write.save("nameAndCity.parquet")
>>> df.select("firstName", "age") .write .save("namesAndAges.json", format="json")
```


Spark MLlib

Apache Spark Components

- Combine SQL, streaming, and complex analytics.




Apache Spark MLlib



- MLlib is developed as part of the Apache Spark project. It thus gets tested and updated with each Spark release.
- MLlib is Spark's scalable machine learning library consisting of common learning algorithms and utilities, including **classification, regression, clustering, collaborative filtering, dimensionality reduction**, and more.
- List of algorithms implemented in MLlib:
<http://spark.apache.org/docs/latest/ml-guide.html>

Spark Documentation: <http://spark.apache.org/docs/latest/ml-guide.html>

 2.4.4

OverviewProgramming Guides▼API Docs▼Deploying▼More▼

MLlib: Main Guide

- Basic statistics
- Data sources
- Pipelines
- Extracting, transforming and selecting features
- Classification and Regression**
- Clustering
- Collaborative filtering
- Frequent Pattern Mining
- Model selection and tuning
- Advanced topics

MLlib: RDD-based API Guide

- Data types
- Basic statistics
- Classification and regression
- Collaborative filtering
- Clustering
- Dimensionality reduction
- Feature extraction and transformation
- Frequent pattern mining
- Evaluation metrics
- PMML model export
- Optimization (developer)


Classification and regression

This page covers algorithms for Classification and Regression. It also includes sections discussing specific classes of algorithms, such as linear methods, trees, and ensembles.

Table of Contents

- Classification
 - Logistic regression
 - Binomial logistic regression**
 - Multinomial logistic regression
 - Decision tree classifier
 - Random forest classifier
 - Gradient-boosted tree classifier
 - Multilayer perceptron classifier
 - Linear Support Vector Machine
 - One-vs-Rest classifier (a.k.a. One-vs-All)
 - Naive Bayes
- Regression
 - Linear regression
 - Generalized linear regression
 - Available families
 - Decision tree regression
 - Random forest regression
 - Gradient-boosted tree regression
 - Survival regression
 - Isotonic regression
- Linear methods
- Decision trees
 - Inputs and Outputs
 - Input Columns
 - Output Columns
- Tree Ensembles
 - Random Forests
 - Inputs and Outputs
 - Input Columns
 - Output Columns (Predictions)
 - Gradient-Boosted Trees (GBTs)
 - Inputs and Outputs
 - Input Columns
 - Output Columns (Predictions)

Spark Documentation: <http://spark.apache.org/docs/latest/ml-guide.html>



OverviewProgramming Guides▼API Docs▼Deploying▼More▼

MLlib: Main Guide

- Basic statistics
- Data sources
- Pipelines
- Extracting, transforming and selecting features
- Classification and Regression**
- Clustering
- Collaborative filtering
- Frequent Pattern Mining
- Model selection and tuning
- Advanced topics

MLlib: RDD-based API Guide

- Data types
- Basic statistics
- Classification and regression
- Collaborative filtering
- Clustering
- Dimensionality reduction
- Feature extraction and transformation
- Frequent pattern mining
- Evaluation metrics
- PMML model export
- Optimization (developer)

Binomial logistic regression

For more background and more details about the implementation of binomial logistic regression, refer to the documentation of [logistic regression in spark.mllib](#).

Examples

The following example shows how to train binomial and multinomial logistic regression models for binary classification with elastic net regularization. `elasticNetParam` corresponds to α and `regParam` corresponds to λ .

ScalaJavaPythonR

```
from pyspark.ml.classification import LogisticRegression

# Load training data
training = spark.read.format("libsvm").load("data/mllib/sample_libsvm_data.txt")

lr = LogisticRegression(maxIter=10, regParam=0.3, elasticNetParam=0.8)

# Fit the model
lrModel = lr.fit(training)

# Print the coefficients and intercept for logistic regression
print("Coefficients: " + str(lrModel.coefficients))
print("Intercept: " + str(lrModel.intercept))

# We can also use the multinomial family for binary classification
mlr = LogisticRegression(maxIter=10, regParam=0.3, elasticNetParam=0.8, family="multinomial")

# Fit the model
mlrModel = mlr.fit(training)

# Print the coefficients and intercepts for logistic regression with multinomial family
print("Multinomial coefficients: " + str(mlrModel.coefficientMatrix))
print("Multinomial intercepts: " + str(mlrModel.interceptVector))
```

More details on parameters can be found in the [Python API documentation](#).

Find full example code at "examples/src/main/python/ml/logistic_regression_with_elastic_net.py" in the Spark repo.

The `spark.ml` implementation of logistic regression also supports extracting a summary of the model over the training set. Note that the predictions and metrics which are stored as `DataFrame` in `LogisticRegressionSummary` are annotated `@transient` and hence only available on the driver.

Preguntas?





































Agenda

- Apache Hadoop
- Apache Spark
- Distribuciones Comerciales de Hadoop
- AWS Elastic Map Reduce (EMR)
- PySpark
- Spark SQL
- Spark MLlib
- Como abrir una cuenta de AWS gratis?

AWS Foundational Services

Compute	Network	Storage	Security & Identity	Applications
 Amazon EC2	 Amazon CloudFront	 Amazon EFS	 Amazon Inspector	 Amazon WorkDocs
 Amazon Elastic Container Registry	 Amazon Route 53	 Amazon S3 Glacier	 AWS Artifact	 Amazon WorkMail
 Amazon Elastic Container Service	 Amazon VPC	 Amazon S3	 AWS Certificate Manager	 Amazon AppStream 2.0
 Amazon Lightsail	 AWS Direct Connect	 AWS Snowball	 AWS CloudHSM	 Amazon WorkSpaces
 AWS Batch	 Elastic Load Balancing	 AWS Storage Gateway	 AWS Directory Service	
 AWS Elastic Beanstalk			 AWS IAM	
 AWS Lambda			 AWS KMS	
			 AWS Organizations	
			 AWS Shield	
			 AWS WAF	

AWS Platform Services

Databases	Analytics	Application Services	Management Tools	Developer Tools	Mobile Services	Internet of Things
 Amazon DynamoDB	 Amazon Athena	 Amazon API Gateway	 Amazon CloudWatch	 AWS CodeBuild	 AWS Mobile Hub	 AWS IoT
 Amazon ElastiCache	 Amazon CloudSearch	 Amazon AppStream 2.0	 AWS CloudFormation	 AWS CodeCommit	 Amazon Cognito	 AWS IoT Greengrass
 Amazon RDS	 Amazon EMR	 Amazon Elastic Transcoder	 AWS CloudTrail	 AWS CodeDeploy	 Amazon Mobile Analytics	
 Amazon Redshift	 Amazon Elasticsearch Service	 Amazon SWF	 AWS Config	 AWS CodePipeline	 Amazon Pinpoint	
	 Amazon Kinesis	 AWS Step Functions	 AWS Managed Services	 AWS X-Ray	 AWS Device Farm	
	 Amazon QuickSight		 AWS OpsWorks	 AWS CodeStar		
			 AWS Service Catalog			
			 AWS Trusted Advisor			

Como abrir una cuenta de AWS gratis?

- AWS Free Tier Account → pide tarjeta de credito
- AWS Educate Account → si eres estudiante, puedes accede sin tarjeta de credito

Preguntas?